Advanced Out-of-Distribution Detection Frameworks for Fine-Grained Plant Disease Diagnosis: A Synthesis of Vision Transformers, Foundation Models, and Parameter-Efficient Adaptation

1. The Operational Imperative: Reliability in Fine-Grained Agricultural Artificial Intelligence

The deployment of automated diagnostic systems within the infrastructure of precision agriculture represents a transformative shift in global food security. Deep learning architectures have achieved superhuman performance in closed-set classification tasks, yet their viability in real-world agricultural environments is fundamentally constrained by their reliability when discerning "known" from "unknown" stimuli.[1] While a model may classify pre-defined pathogen categories with high accuracy, the operational reality of agricultural fields introduces a vast, stochastic array of Out-of-Distribution (OOD) inputs—ranging from novel pathogen strains and abiotic environmental stressors to non-plant artifacts, sensor anomalies, and variable lighting conditions.[1]

The inability of standard classifiers to robustly reject these anomalies constitutes a critical safety vulnerability. In the high-stakes context of phytopathology, a high-confidence misclassification of a novel, virulent blight as a benign nutrient deficiency can lead to unchecked spread and catastrophic crop loss.[1] This report provides a comprehensive technical analysis of the state-of-the-art methodologies for OOD detection, specifically tailored to the architectural paradigm of Vision Transformers (ViTs) and the domain constraints of Fine-Grained Visual Categorization (FGVC).

1.1 The Fine-Grained Visual Categorization (FGVC) Challenge in Pathology

Plant disease diagnosis is a quintessential Fine-Grained Visual Categorization (FGVC) problem. Unlike generic object recognition, which might involve discriminating between disparate categories such as "vehicle" and "animal," FGVC requires the discrimination of subclasses that share a high degree of structural similarity.[1] In the context of phytopathology, a "healthy" tomato leaf, a leaf with "Early Blight" (*Alternaria solani*), and a leaf with "Septoria Leaf Spot" (*Septoria lycopersici*) share the same global geometry, color palette, and biological morphology. The discriminative features are often minute, localized textural anomalies—concentric rings in a necrotic spot versus water-soaked lesions—that occupy a negligible fraction of the image pixels.[1]

This high intra-class variance (due to varying growth stages, lighting, and leaf angles) and low inter-class variance creates a perilous landscape for OOD detection. A standard Convolutional Neural Network (CNN) or Vision Transformer (ViT) trained on a closed set of 10 diseases will partition the high-dimensional feature space into 10 regions.[1] When presented with an 11th, unknown disease (OOD), the model forces the input into one of the existing partitions. Because the global structure (a leaf) is In-Distribution (ID), the model often assigns high confidence to the prediction based on shared background features, a phenomenon known as the "overconfidence" problem or "high-confidence fool".[1]

1.2 The Shift from CNNs to Vision Transformers

Historically, CNNs dominated this field, leveraging their inductive bias for local texture to identify lesions. However, the field is rapidly shifting toward Vision Transformers (ViTs).[1] ViTs, which process images as sequences of patch tokens, utilize self-attention mechanisms capable of modeling long-range dependencies. This is theoretically advantageous for plant disease detection, where the spatial distribution of lesions (e.g., scattered spots vs. marginal necrosis) is diagnostic.[1]

However, ViTs introduce new complexities for OOD detection that were less prevalent in CNN architectures:

- **Feature Uniformity:** The Layer Normalization (LayerNorm) inherent in ViTs tends to produce feature representations that are more uniform in magnitude and distribution compared to the unnormalized activations of CNNs. This uniformity can obscure the "surprise" signal typically used to detect anomalies.[1]
- **Shape Bias:** ViTs exhibit a stronger "shape bias" compared to the "texture bias" of CNNs. While this improves robustness to occlusion, it can make the model less sensitive to the textural anomalies that distinguish different pathologies, potentially clustering distinct diseases closer together in the latent space.[1]
- **OOD Behavior:** Empirical studies suggest that standard post-hoc detection methods optimized for CNNs (like ODIN or React) often degrade in performance when applied directly to ViTs without modification.[1]

1.3 The Scope of Analysis

This report synthesizes diverse research streams to propose a unified, robust framework for agricultural OOD detection. We dissect specific mechanisms including:

1. **Geometric Methods:** The Mahalanobis Distance and its evolutions (RMD, Mahalanobis++), which leverage the covariance structure of the feature space.[1]
2. **Foundation Model Strategies:** The specific utility of DINOv2 and the trade-off between

linear probing and non-parametric nearest-neighbor evaluations.[1]

3. **Adaptation-Based Detection:** The novel exploitation of LoRA parameters as uncertainty sensors, including unmerged embeddings, selective low-rank approximation (SeTAR), and boxed abstraction monitors (LoRA-BAM).[1]
4. **Continual Learning Integration:** Frameworks like BUILD, TPL, and RCL that bridge the gap between static OOD detection and dynamic, lifelong learning in non-stationary agricultural environments.[17]

2. Geometric Approaches: Mahalanobis Distance and Its Evolutions

The geometric interpretation of neural feature spaces provides the most mathematically grounded framework for OOD detection. The core hypothesis is that a well-trained network maps ID data to a union of compact, low-dimensional manifolds (approximated as Gaussians), and OOD data falls into the low-density regions between or outside these manifolds.[1]

2.1 Theoretical Foundations of Mahalanobis Distance (MD)

The Mahalanobis Distance (MD) is a generalized distance metric that accounts for the correlations between variables in a dataset. Unlike the Euclidean distance, which assumes that features are uncorrelated and have unit variance (isotropic), MD "whitens" the data based on the empirical covariance matrix.[1] Formally, consider a pre-trained feature extractor $f(x)$ that maps an input image to a feature vector $z \in \mathbb{R}^d$ in the penultimate layer. We model the distribution of features for each class $c \in \{1, ..., C\}$ as a multivariate Gaussian distribution $\mathcal{N}(\mu_c, \Sigma)$.[1]

The class-conditional mean is estimated as:

$$\hat{\mu}_c = \frac{1}{N_c} \sum_{i:y_i=c} f(x_i)$$

To ensure robust estimation, especially in high-dimensional spaces where the number of samples per class (N_c) might be small relative to the dimensionality (d), we typically assume a **tied covariance matrix** $\hat{\Sigma}$ shared across all classes. This is calculated by pooling the sample covariances [1]:

$$\hat{\Sigma} = \frac{1}{N} \sum_{c=1}^{C} \sum_{i:y_i=c} (f(x_i) - \hat{\mu}_c)(f(x_i) - \hat{\mu}_c)^T$$

The OOD score for a test sample is defined as the minimum squared distance to any class centroid, scaled by the inverse covariance:

$$M(x) = \min_c (f(x) - \hat{\mu}_c)^T \hat{\Sigma}^{-1} (f(x) - \hat{\mu}_c)$$

Why MD Outperforms Softmax: The Softmax function in the final layer is a projection that compresses the high-dimensional feature vector into a probability simplex. This compression inevitably results in information loss. The Softmax logits effectively represent the distance to the decision boundary (hyperplane), not the distance to the class prototype. A sample can be extremely far from the centroid (an outlier) but still be far from the decision boundary (high confidence), leading to the "high-confidence fool" problem. MD, by operating in the feature space, captures the distance to the density mode, providing a direct measure of "typicality".[1]

2.2 The "Simple Fix": Relative Mahalanobis Distance (RMD)

Despite the theoretical elegance of MD, it suffers from a critical failure mode in "Near-OOD" scenarios—situations where the OOD samples share significant semantic overlap with the ID data. This is precisely the case in plant disease detection, where a "Tomato Yellow Leaf Curl" image (OOD) looks structurally identical to a "Tomato Mosaic Virus" image (ID) except for specific coloration patterns.[1]

The Background Confounding Problem: In deep neural networks, the magnitude (norm) of the feature vector carries significant information about the "background" content of the image. For instance, the presence of a leaf shape and green pixels triggers a baseline level of activation across the network filters. This "background signal" contributes to the total distance $M(x)$. When both ID and Near-OOD samples share this background, their $M(x)$ scores become indistinguishable, as the shared background distance dominates the subtle class-specific distance.[1]

The RMD Methodology: Ren et al. proposed the Relative Mahalanobis Distance (RMD) to isolate the class-specific signal. The method assumes that the feature vector is composed of a class-specific component and a class-agnostic (background) component. To estimate the background contribution, RMD fits a second Gaussian distribution $\mathcal{N}(\mu_0, \Sigma_0)$ to the entire training set, ignoring class labels. The RMD score is the difference between the class-specific distance and the background distance [1]:

$$RMD(x) = M(x) - M_0(x)$$

where $M_0(x) = (f(x) - \hat{\mu}_0)^T \hat{\Sigma}_0^{-1} (f(x) - \hat{\mu}_0)$. This formulation is mathematically equivalent to a log-likelihood ratio test between a class-conditional model and a background model. By subtracting the background term, RMD cancels out the common factors (e.g., "it is a leaf"). Empirical validation on genomics and FGVC benchmarks has shown RMD to improve AUROC by nearly 15.8 points compared to standard MD in scenarios with high semantic overlap.[1]

2.3 Mahalanobis++: Addressing Feature Norm Instability

While RMD addresses semantic overlap, recent research has highlighted a structural instability in MD related to the statistical properties of ViT features. The study "Mahalanobis++" identifies that the L_2 norms of feature vectors in modern pre-trained models are often heavy-tailed and not normally distributed.[1]

The Norm-Variance Correlation: In many models, there is a strong correlation between the norm of the feature vector $||z||_2$ and the prediction confidence. OOD samples often result in feature vectors with significantly smaller (or larger) norms than ID samples. However, the standard MD calculation allows the covariance matrix to be dominated by the directions of high variance, which are often just the directions of magnitude change. This creates "blind spots" where OOD samples with typical norms but atypical angles are not detected.[1]

The Normalization Solution: Mahalanobis++ proposes a preprocessing step: L_2 -normalization of the features before computing the Gaussian statistics.

$$z_{norm} = \frac{z}{||z||_2}$$

By projecting all features onto the unit hypersphere, the method eliminates the variance due to magnitude. The MD then becomes purely a measure of angular distance (cosine similarity) scaled by the angular spread of the class cluster. Experimental results on 44 different models (including ViTs and ConvNeXts) showed that this normalization consistently improved OOD detection performance, reducing the False Positive Rate at 95% True Positive Rate (FPR95) by an average of 9.6% compared to the conventional Mahalanobis score.[11] For Vision Transformers in agriculture, which often process images with variable background clutter (leading to variable feature energy), Mahalanobis++ provides a crucial stabilization mechanism.[10]

3. Foundation Model Strategies: DINOv2 and the OOD Landscape

The paradigm of training models from scratch (supervised learning) is being superseded by the use of Foundation Models. DINOv2 (Discriminative Self-supervised Learning) represents the state-of-the-art in this domain, offering feature spaces that are remarkably robust for OOD detection.[1]

3.1 DINOv2: Architecture and Pre-training Signals

DINOv2 employs a student-teacher architecture trained with a combination of DINO (self-distillation) and iBOT (Masked Image Modeling) losses.[1]

- **DINO Loss:** Encourages the student network to match the teacher's output on different augmented views of the same image (global consistency).
- **iBOT Loss:** Forces the student to reconstruct masked patches of the image based on the visible context (local consistency).

This dual objective is critical for fine-grained plant disease tasks. Unlike CLIP, which aligns images to text captions (often losing fine-grained visual details), DINOv2 is purely visual and learns features that capture subtle spatial patterns required for depth and fine-grained segmentation.[1] The iBOT loss compels the model to understand local textures (e.g., "this green patch implies the neighboring patch should be green," or "this yellow halo implies a fungal center").

3.2 The Linear Probe vs. Nearest Neighbor (k-NN) Debate

When adapting a foundation model like DINOv2, a critical choice arises between using a Linear Probe or Non-Parametric Nearest Neighbor (k-NN) evaluation.[1]

Linear Probing (The Bottleneck Effect): Linear probing involves freezing the DINOv2 backbone and training a linear layer on the labeled ID dataset. While effective for classification, it can degrade OOD detection. The linear projection simplifies the complex, high-dimensional manifold of the DINOv2 features into a lower-dimensional decision space. OOD samples that lie far from the ID manifold in the feature space might project onto the "high confidence" side of the linear boundary in the decision space—a phenomenon known as "feature collapse".[1]

Nearest Neighbor (k-NN - Manifold Preservation): The k-NN approach stores feature embeddings of the training set. For a test sample, the OOD score is the distance to the k-th nearest neighbor. k-NN operates directly on the native DINOv2 manifold. Because DINOv2 is trained to cluster visually similar images, ID samples naturally cluster tightly, while OOD samples fall into sparse regions. Empirical evidence suggests that DINOv2 + k-NN consistently outperforms linear probes on challenging benchmarks like iNaturalist and NINCO, making it ideal for agricultural scenarios where the "long tail" of rare diseases makes training robust classifiers difficult.[1]

3.3 Visual Prompt Tuning (VIPAMIN)

Recent advancements in parameter-efficient tuning have introduced **VIPAMIN** (Visual Prompt Initialization via Embedding Selection and Subspace Expansion) as a superior alternative to standard Visual Prompt Tuning (VPT) for OOD robustness.[25] Unlike standard VPT, which initializes prompts randomly, VIPAMIN employs two specialized modules:

1. **Matching Module:** Aligns prompts with semantically informative input tokens from the downstream task (e.g., specific disease lesions). This addresses the issue of uniform attention in standard ViTs, promoting specialization over meaningful local regions.[25]
2. **Orthogonalizing Module:** Prevents representational collapse by projecting prompts away from the pre-trained embedding subspace. This facilitates the injection of novel, task-specific information required to distinguish fine-grained classes.[25]

Experiments on CIFAR-100 corrupted benchmarks show that VIPAMIN improves OOD generalization accuracy by 4.0% over standard prompt tuning, making it a critical component for robust agricultural models operating in variable field conditions.[26]

4. Adaptation-Based Detection: LoRA and SeTAR

While DINOv2 provides general features, specialized crop disease tasks benefit from Low-Rank Adaptation (LoRA). Modern research has moved beyond simple merging to exploit LoRA modules as active uncertainty sensors.[1]

4.1 Unmerged LoRA Embeddings (LoRA-MD)

A significant breakthrough involves using the activations of the LoRA branch (BAx) before they are merged into the pre-trained weights.[1] The LoRA embedding $E_{LoRA}(x)$ is defined as the concatenation of intermediate activations $A_j x$ for all L layers. The LoRA parameters A and B are trained solely on the ID dataset. Therefore, the term BAx represents the "adaptation signal"—the specific features required to process the ID data that were missing from the pre-trained backbone.[1]

In fine-grained scenarios where standard MD on last-layer activations fails (AUROC ~0.4), MD using unmerged LoRA embeddings can reach AUROCs as high as 0.890.[1] This is because the adapter specifically captures the "delta" required to process ID diseases (e.g., recognizing *Septoria* spots), while effectively ignoring general background features that trigger the backbone.

4.2 Selective Low-Rank Approximation (SeTAR)

SeTAR is a novel, training-free OOD detection method that leverages selective low-rank approximation of weight matrices.[1] Standard adaptation often introduces noise via minor singular components. SeTAR mitigates this through a sophisticated greedy search algorithm:

- **Greedy Search Algorithm:** The algorithm iterates through the network layers (typically top-to-bottom, image-to-text encoders in CLIP-like models) to identify the optimal rank

reduction ratio for each weight matrix.
- **Metric:** It evaluates potential configurations using the **LoCoOp loss** on an ID validation set. The LoCoOp loss utilizes pseudo-OOD features generated from background patches (ID-irrelevant nuisances) to ensure the retained components maximize the separation between ID objects and background noise.[30]
- **Fine-Tuning Extension (SeTAR+FT):** This extension fine-tunes the "minor" singular components of weight matrices while freezing the "major" ones. This counter-intuitive strategy stabilizes the OOD-robust manifold by allowing the model to adapt to specific ID nuances without disrupting the core feature structures learned during pre-training.[15]

On ImageNet benchmarks, SeTAR+FT has been shown to reduce false positive rates by up to 18.95% compared to zero-shot baselines.[15]

4.3 LoRA-BAM: Boxed Abstraction Monitors

LoRA-BAM introduces an interpretable, geometric filtering layer directly over the LoRA modules.[1] Unlike statistical distance methods, LoRA-BAM constructs explicit boundaries:

1. **Feature Extraction:** Activations are extracted from the LoRA adaptation layers.
2. **Clustering:** The algorithm applies **k-means clustering** to these vectors to identify dense regions of ID behavior.[14]
3. **Box Construction:** It constructs "boxed abstractions"—conservative, axis-aligned geometric bounds—around these clusters. The bounds are often enlarged based on feature variance to accommodate natural ID variations (e.g., lighting changes or paraphrasing in NLP).[14]
4. **Inference:** A query is flagged as OOD if its feature vector falls outside the union of all defined boxes.

LoRA-BAM is particularly effective for rejecting "Far-OOD" samples that are semantically distinct from the training data. Empirical results show it can reject up to 95% of Far-OOD queries while retaining nearly all legitimate ID samples, reducing hallucination errors by 88% in open-world detection benchmarks.[32]

5. The Continual Learning Frontier: Evolving with the Environment

In agricultural practice, the environment is non-stationary. New diseases emerge, and abiotic stress patterns shift with climate change. Models must therefore employ **Continual Learning (CL)** to adapt without catastrophic forgetting. Crucially, successful CL relies on robust OOD detection to identify *when* a new class has appeared.[37]

5.1 The Theoretical Nexus: Task-ID Prediction

The Class Incremental Learning (CIL) problem can be decomposed into Within-Task Prediction (WP) and Task-ID Prediction (TP).[37] TP is fundamentally an OOD detection problem: determining if an input belongs to the current task or a new, unseen distribution. Research indicates a strong linear correlation (Pearson $r \approx 0.976$) between OOD detection accuracy and CIL performance.[39]

5.2 Frameworks for Continual OOD Detection

BUILD (Buffer-free Incremental Learning with OOD Detection): Traditional CL methods use memory buffers (Replay) to store old data, which raises privacy and scalability concerns. **BUILD** is a buffer-free framework that integrates a pre-trained ViT with hard attention masks and post-hoc OOD detectors.[17] It uses hard attention to "lock" parameters associated with old tasks, preventing forgetting, while using activation-based OOD detectors (like Mahalanobis++) to assign incoming samples to the correct task-specific adapter. Evaluations on CIFAR-10 show BUILD achieves superior stability and lower computational cost than buffer-based methods like MORE.[17]

RCL (Reliable Continual Learning): RCL addresses the "Unified Failure Detection" problem, aiming to detect both misclassified ID samples and OOD samples. It employs a **Weight Space Interpolation (WSI)** strategy, ensembling models along the fine-tuning trajectory to capture diverse uncertainty estimates. This is critical for differentiating between "hard" ID samples (e.g., a confusingly lit leaf) and true OOD samples (e.g., a completely new virus).[19]

TPL (Task-id Prediction via Likelihood Ratio): TPL improves TP by leveraging replay data to estimate the likelihood ratio between the current task and past tasks. Instead of a simple threshold, it calculates the probability ratio $P(x|Task_{current})/P(x|Task_{past})$, offering a more principled decision boundary than standard OOD scores.[18]

SCALE and ISH (Intermediate Tensor Shaping): To enhance OOD detection during the learning process, the **SCALE** framework introduces **Intermediate Tensor Shaping (ISH)**. Unlike activation pruning, which can discard useful information, ISH *scales* intermediate feature maps during training to emphasize ID characteristics. This training-time enhancement improves the separability of ID and OOD data in the latent space, achieving significant AUROC gains (+1.85% on Near-OOD ImageNet benchmarks) with minimal computational overhead.[43]

5.3 Agricultural Application: The "Discover-and-Learn" Cycle

In applied settings like strawberry disease monitoring, these frameworks enable a "Discover-and-Learn" cycle. An initial model differentiates known diseases from unknown anomalies (OOD). These unknowns are clustered (using methods like OpenMatch or two-head networks) and assigned temporary pseudo-labels. In a subsequent training round, the model incorporates these new clusters as distinct classes, progressively expanding its knowledge base without human intervention.[48] This approach has been successfully validated in wheat disease classification, maintaining 98% accuracy while adapting to new stress markers.[2]

6. Synthesis and Recommendation for Agricultural Deployment

For the specific challenge of Fine-Grained Plant Disease Detection, relying on a single detection method is insufficient. The diverse nature of anomalies—from sensor noise (Far-OOD) to visually similar viral strains (Near-OOD)—requires a multi-layered architectural framework.

6.1 The Integrated Framework

Component	**Technology**	**Implementation Detail**	**Primary Function**
Backbone	**DINOv2 (ViT-L/14)**	Pre-trained with DINO + iBOT losses.	Captures fine-grained texture/shape features critical for lesions.
Adaptation	**SeTAR+FT**	Selective low-rank fine-tuning (freezing major, tuning minor components).	Adapts to specific crop/disease while preserving OOD-robust manifold stability.
Normalization	**Mahalanobis++**	L_2-normalization of feature vectors before statistical analysis.	Stabilizes angular distance metrics against variable image contrast/lighting.
Core Detection	**LoRA-MD**	Calculate MD on concatenated, unmerged adapter activations (BAx).	Detects "Near-OOD" diseases by measuring "adaptation surprise."

Filtering	**LoRA-BAM**	k-means clustering + boxed abstraction on LoRA layers.	Rejects "Far-OOD" artifacts (soil, tools) with interpretable geometric bounds.
Lifecycle	**BUILD / RCL**	Buffer-free incremental learning with hard attention masks.	Enables continuous adaptation to new strains without catastrophic forgetting.
Validation	**Grad-CAM / MECAM**	Multi-exit attention mapping.	Visual sanity check to ensure focus is on pathology, not background.

6.2 Conclusion

The evolution of Out-of-Distribution detection for agricultural Vision Transformers signifies a transition from global probabilistic tail-modeling (OpenMax) to precise local geometric manifold analysis. The limitations of classical Extreme Value Theory in high-dimensional latent spaces have been effectively addressed through the development of Relative Mahalanobis Distance and feature normalization techniques like Mahalanobis++.

By integrating OOD awareness directly into the adaptation layer—through unmerged LoRA embeddings, selective rank approximation (SeTAR), and boxed abstraction monitors (LoRA-BAM)—modern frameworks achieve the high sensitivity required to distinguish between semantically proximate diseases. Furthermore, the unification of OOD detection with buffer-free continual learning (BUILD, RCL) allows these systems to evolve alongside the dynamic ecosystems they monitor. This multi-layered approach provides the high-fidelity reliability and interpretability required to bridge the gap between laboratory success and robust, open-world agricultural deployment. The future of agricultural AI lies not just in recognizing the known, but in the intelligent, continuous discovery and management of the unknown.

Alıntılanan çalışmalar

1. Advanced Out-of-Distribution Detection Frameworks for Fine-Grained Plant Disease Diagnosis_ A Synthesis of Vision Transformers, Foundation Models, and Parameter-Efficient Adaptation.pdf
2. Regional feature purification contrastive learning for wheat biotic stress detection

- PMC, erişim tarihi Şubat 4, 2026, https://pmc.ncbi.nlm.nih.gov/articles/PMC12558983/
3. Unsupervised Incremental Learning Framework for Online Fault Diagnosis | Industrial & Engineering Chemistry Research - ACS Publications, erişim tarihi Şubat 4, 2026, https://pubs.acs.org/doi/10.1021/acs.iecr.5c00949?ai=550&mi=rpypxk&af=R
4. TOM-SSL: Tomato Disease Recognition Using Pseudo-Labelling-Based Semi-Supervised Learning - MDPI, erişim tarihi Şubat 4, 2026, https://www.mdpi.com/2624-7402/7/8/248
5. Two-head network for Out-of-Distribution (OoD) detection. - ResearchGate, erişim tarihi Şubat 4, 2026, https://www.researchgate.net/figure/Two-head-network-for-Out-of-Distribution-OoD-detection_fig3_363600523
6. DISSECTING MAHALANOBIS: HOW FEATURE GEOMETRY AND NORMALIZATION SHAPE OOD DETECTION - OpenReview, erişim tarihi Şubat 4, 2026, https://openreview.net/pdf/18dccafdbd8986b5ee0bb5b29435ec0be3159508.pdf
7. Exploring the Limits of Out-of-Distribution Detection - NeurIPS, erişim tarihi Şubat 4, 2026, https://proceedings.neurips.cc/paper/2021/hash/3941c4358616274ac2436eacf67fae05-Abstract.html
8. Combining pre-trained Vision Transformers and CIDER for Out Of Domain Detection - arXiv, erişim tarihi Şubat 4, 2026, https://arxiv.org/abs/2309.03047
9. Scaling Vision Transformers to 22 Billion Parameters - Proceedings of Machine Learning Research, erişim tarihi Şubat 4, 2026, https://proceedings.mlr.press/v202/dehghani23a/dehghani23a.pdf
10. mueller-mp/maha-norm - GitHub, erişim tarihi Şubat 4, 2026, https://github.com/mueller-mp/maha-norm
11. Mahalanobis++: Improving OOD Detection via Feature Normalization - arXiv, erişim tarihi Şubat 4, 2026, https://arxiv.org/html/2505.18032v1
12. (PDF) DINOv2 Rocks Geological Image Analysis: Classification, Segmentation, and Interpretability - ResearchGate, erişim tarihi Şubat 4, 2026, https://www.researchgate.net/publication/382560010_DINOv2_Rocks_Geological_Image_Analysis_Classification_Segmentation_and_Interpretability
13. DINOv2 Rocks Geological Image Analysis: Classification, Segmentation, and Interpretability, erişim tarihi Şubat 4, 2026, https://arxiv.org/html/2407.18100v3
14. Input Filtering for Fine-tuned LLMs via Boxed Abstraction Monitors over LoRA Layers - arXiv, erişim tarihi Şubat 4, 2026, https://arxiv.org/pdf/2506.00998
15. SeTAR, erişim tarihi Şubat 4, 2026, https://setar-ood.github.io/
16. Etienne Salimbeni | Semantic Scholar, erişim tarihi Şubat 4, 2026, https://www.semanticscholar.org/author/Etienne-Salimbeni/2296462889
17. BUILD: Buffer-free Incremental Learning with OOD Detection for the Wild - OpenReview, erişim tarihi Şubat 4, 2026, https://openreview.net/pdf?id=WaGj1QfgCH
18. Class-Incremental Learning (CIL) - Emergent Mind, erişim tarihi Şubat 4, 2026, https://www.emergentmind.com/topics/class-incremental-learning-cil

19. RCL: Reliable Continual Learning for Unified Failure Detection - ResearchGate, erişim tarihi Şubat 4, 2026, https://www.researchgate.net/publication/384144015_RCL_Reliable_Continual_Learning_for_Unified_Failure_Detection
20. Dissecting Mahalanobis: How Feature Geometry and Normalization Shape OOD Detection, erişim tarihi Şubat 4, 2026, https://arxiv.org/html/2510.15202v2
21. ICML Poster Mahalanobis++: Improving OOD Detection via Feature Normalization, erişim tarihi Şubat 4, 2026, https://icml.cc/virtual/2025/poster/43649
22. Mahalanobis++: Improving OOD Detection via Feature Normalization - GitHub, erişim tarihi Şubat 4, 2026, https://raw.githubusercontent.com/mlresearch/v267/main/assets/muller25a/muller25a.pdf
23. Towards a General-Purpose Foundation Model for Computational Pathology - PMC, erişim tarihi Şubat 4, 2026, https://pmc.ncbi.nlm.nih.gov/articles/PMC11403354/
24. Evaluating the generalizability of a panorama-point cloud encoder trained without supervision - kth .diva, erişim tarihi Şubat 4, 2026, https://kth.diva-portal.org/smash/get/diva2:2010937/FULLTEXT01.pdf
25. VIPAMIN: Visual Prompt Initialization via Embedding Selection and Subspace Expansion, erişim tarihi Şubat 4, 2026, https://arxiv.org/html/2510.16446v1
26. VIPAMIN: Visual Prompt Initialization via Embedding Selection and Subspace Expansion - OpenReview, erişim tarihi Şubat 4, 2026, https://openreview.net/pdf?id=o0AASFieVc
27. How Can LoRA parameters improve the detection of Near-OOD data? | Mostafa ElAraby, erişim tarihi Şubat 4, 2026, https://www.mostafaelaraby.com/paper%20review/2025/11/03/LoRA-parameters-improves-OOD-detection/
28. BEYOND FINE-TUNING: LORA MODULES BOOST NEAR- OOD DETECTION AND LLM SECURITY - OpenReview, erişim tarihi Şubat 4, 2026, https://openreview.net/pdf?id=H7Q5hHcvZE
29. SeTAR: Out-of-Distribution Detection with Selective Low-Rank Approximation - arXiv, erişim tarihi Şubat 4, 2026, https://arxiv.org/html/2406.12629v1
30. SeTAR: Out-of-Distribution Detection with Selective Low-Rank Approximation - OpenReview, erişim tarihi Şubat 4, 2026, https://openreview.net/pdf?id=65UoJ0z7Kp
31. SeTAR: Out-of-Distribution Detection with Selective Low-Rank Approximation - arXiv, erişim tarihi Şubat 4, 2026, https://arxiv.org/html/2406.12629v4
32. BAM: Box Abstraction Monitors for Real-time OoD Detection in Object Detection, erişim tarihi Şubat 4, 2026, https://www.semanticscholar.org/paper/BAM%3A-Box-Abstraction-Monitors-for-Real-time-OoD-in-Wu-He/34d8cb340082f0662939eb4a02363ac3c1d86c34
33. BAM: Box Abstraction Monitors for Real-time OoD Detection in Object Detection | Request PDF - ResearchGate, erişim tarihi Şubat 4, 2026, https://www.researchgate.net/publication/387427660_BAM_Box_Abstraction_Monitors_for_Real-time_OoD_Detection_in_Object_Detection

34. Daily Papers - Hugging Face, erişim tarihi Şubat 4, 2026, https://huggingface.co/papers?q=LoRA
35. Input Filtering for Fine-tuned LLMs via Boxed Abstraction Monitors over LoRA Layers - arXiv, erişim tarihi Şubat 4, 2026, https://arxiv.org/html/2506.00998v1
36. LLM Response Filtering Techniques - Emergent Mind, erişim tarihi Şubat 4, 2026, https://www.emergentmind.com/topics/response-filtering-in-llms
37. A Theoretical Study on Solving Continual Learning - Computer Science, erişim tarihi Şubat 4, 2026, https://www.cs.uic.edu/~liub/publications/NeurIPS-2022-camera-ready.pdf
38. Unifying Continual Learning, OOD - Detection & Open World Learning - Computer Science - University of Illinois Chicago, erişim tarihi Şubat 4, 2026, https://www.cs.uic.edu/~liub/Unify-CIL-OOD-OWL.pdf
39. Class Incremental Learning via Likelihood Ratio Based Task Prediction - Liner, erişim tarihi Şubat 4, 2026, https://liner.com/review/class-incremental-learning-via-likelihood-ratio-based-task-prediction
40. Buffer-free class-incremental learning with out-of-distribution detection - UniCA IRIS, erişim tarihi Şubat 4, 2026, https://iris.unica.it/retrieve/80a50a43-7066-45c2-8c77-2e3d4cb510f7/Buffer-free%20class-incremental%20learning%20with%20out-of-distribution%20detection_2026.pdf
41. RCL: Reliable Continual Learning for Unified Failure Detection - CVF Open Access, erişim tarihi Şubat 4, 2026, https://openaccess.thecvf.com/content/CVPR2024/papers/Zhu_RCL_Reliable_Continual_Learning_for_Unified_Failure_Detection_CVPR_2024_paper.pdf
42. Class Incremental Learning via Likelihood Ratio Based Task Prediction - arXiv, erişim tarihi Şubat 4, 2026, https://arxiv.org/html/2309.15048v4
43. Domain Feature Collapse: Implications for Out-of-Distribution Detection and Solutions, erişim tarihi Şubat 4, 2026, https://arxiv.org/html/2512.04034
44. (PDF) ScaleISH: Scaling for Training-Time and Post-hoc Out-of-distribution Detection Enhancement - ResearchGate, erişim tarihi Şubat 4, 2026, https://www.researchgate.net/publication/382083296_ScaleISH_Scaling_for_Training-Time_and_Post-hoc_Out-of-distribution_Detection_Enhancement
45. SCALING FOR TRAINING-TIME AND POST-HOC OUT- OF-DISTRIBUTION DETECTION ENHANCEMENT - ICLR Proceedings, erişim tarihi Şubat 4, 2026, https://proceedings.iclr.cc/paper_files/paper/2024/file/46d943bc6a15a57c923829efc0db7c7a-Paper-Conference.pdf
46. DisCoPatch: Batch Statistics Are All You Need For OOD Detection, But Only If You Can Trust Them - arXiv, erişim tarihi Şubat 4, 2026, https://arxiv.org/html/2501.08005v2
47. arXiv:2310.00227v1 [cs.CV] 30 Sep 2023, erişim tarihi Şubat 4, 2026, https://arxiv.org/pdf/2310.00227
48. A New Deep Learning-based Dynamic Paradigm Towards Open-World Plant Disease Detection - Frontiers, erişim tarihi Şubat 4, 2026, https://www.frontiersin.org/journals/plant-science/articles/10.3389/fpls.2023.1243

822/full
49. A New Deep Learning-based Dynamic Paradigm Towards Open-World Plant Disease Detection - NIH, erişim tarihi Şubat 4, 2026, https://pmc.ncbi.nlm.nih.gov/articles/PMC10577201/
50. Detection of unknown strawberry diseases based on OpenMatch and two-head network for continual learning - NIH, erişim tarihi Şubat 4, 2026, https://pmc.ncbi.nlm.nih.gov/articles/PMC9520169/
51. Detection of unknown strawberry diseases based on OpenMatch and two-head network for continual learning - PubMed, erişim tarihi Şubat 4, 2026, https://pubmed.ncbi.nlm.nih.gov/36186017/